

Probabilistic Algorithms for Aerospace Autonomy  
ASEN 6519 - Homework 2  
Maximum Likelihood Parameter Estimation

Carl Stahoviak  
Carl Mueller

March 13, 2019

## Contents

<b>Problem 1 - Generative Classification</b>	<b>2</b>
<b>Problem 2 - Baum-Welch for Supervisory Operator HMM</b>	<b>5</b>
<b>Appendix - MATLAB Code</b>	<b>7</b>

## Problem 1 - Generative Classification

Estimate the unknown model parameters for model (a) via maximum likelihood and model (b) using ML or EM if necessary.

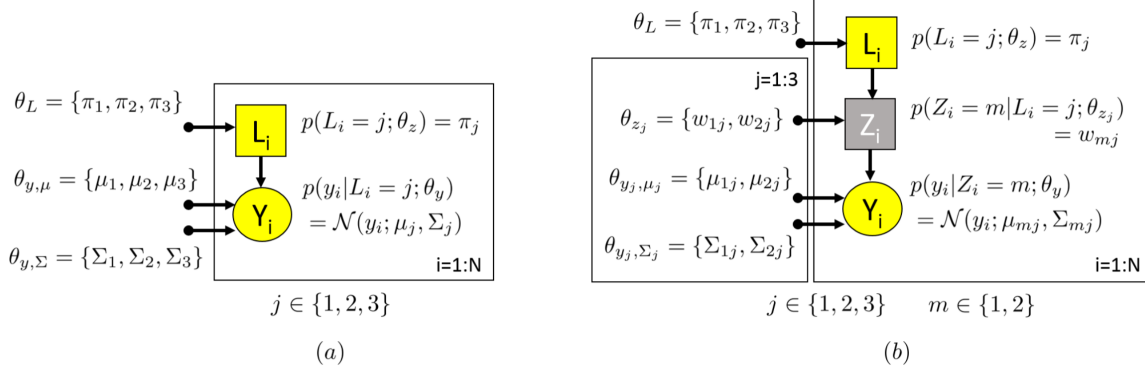


Figure 1: CLL is known for model (a) whereas only the ICLL is known for (b).

### Maximum Likelihood for Model (a)

Below outlines the justification for the chosen implementation for ML the model:

$$\begin{aligned}
 L(y_{1:N}; \theta_j, \theta_y) &= \prod_{i=1}^N p(L_i = j; \theta_L) p(y_i | L_i = j; \theta_y) \\
 l(y_{1:N}; \theta_j, \theta_y) &= \log \left[ \prod_{i=1}^N p(L_i = j; \theta_L) p(y_i | L_i = j; \theta_y) \right] \\
 l(y_{1:N}; \theta_j, \theta_y) &= \log[p(L_i = j; \theta_L)] + \sum_{i=1}^N \log[p(y_i | L_i = j; \theta_y)] \\
 l(y_{1:N}; \theta_j, \theta_y) &= \log \left[ \prod_{j=1}^3 p(\pi_j^{I(L_i=j)}) \right] + \sum_{i=1}^N \log \left[ \prod_{j=1}^3 \mathcal{N}(y_i; \mu_j, \Sigma_j)^{I(L_i=j)} \right] \\
 l(y_{1:N}; \theta_j, \theta_y) &= \sum_{j=1}^3 I(L_i = j) \log[\pi_j] + \sum_{i=1}^N I(L_i = j) \sum_{j=1}^3 \log[\mathcal{N}(y_i; \mu_j, \Sigma_j)] \\
 l(y_{1:N}; \theta_j, \theta_y) &= \sum_{j=1}^3 I(L_i = j) \log[\pi_j] + \sum_{i=1}^N I(L_i = j) \sum_{j=1}^3 \log \left[ \frac{-\frac{1}{2}(y - \mu_j)^T \Sigma_j^{-1} (y - \mu_j)}{\sqrt{(2\pi)^d |\Sigma_j|}} \right]
 \end{aligned}$$

When maximizing the log-likelihood for any of the individual parameters  $\theta_j$ , the indicator function  $I(L_i = j)$  filters out all the data such that we can effectively maximize  $\mu_j$  and  $\Sigma_j$  using data labeled as class  $j$ . Maximizing the partial derivative w.r.t  $\pi_j$  is unnecessary as  $\pi_j$  is simply constant.

Maximizing the partial derivative w.r.t  $\mu_j$  sifts all data out that is not labeled as j, leaving:

$$\frac{\partial I(L_i = j) \log[p(L_i = j; \theta_L)j] + \sum_{i=1}^N I(L_i = j) \sum_{j=1}^3 \log\left[\frac{-\frac{1}{2}(y - \mu_j)^T \Sigma^{-1}(y - \mu_j)}{\sqrt{(2\pi)^d |\Sigma_j|}}\right]}{\partial \mu_j} = 0$$

$$\hat{\mu}_j = \frac{1}{N_j} \sum_j y_j$$

Maximizing the partial derivative w.r.t  $\Sigma$  and setting to zero:

$$\frac{\partial I(L_i = j) \log[p(L_i = j; \theta_L)j] + \sum_{i=1}^N I(L_i = j) \sum_{j=1}^3 \log\left[\frac{-\frac{1}{2}(y - \mu_j)^T \Sigma^{-1}(y - \mu_j)}{\sqrt{(2\pi)^d |\Sigma_j|}}\right]}{\partial \Sigma_j} = 0$$

$$\hat{\Sigma} = \frac{1}{N_j} \sum_j (y_j - \mu_j)(y_j - \mu_j)^T$$

**Estimated Parameters for Model (a):**

$$\mu_1 = \begin{bmatrix} -1.3650 \\ -2.6720 \end{bmatrix}, \Sigma_1 = \begin{bmatrix} 2.5780 & 0.5999 \\ 0.5999 & 2.2415 \end{bmatrix}, \mu_2 = \begin{bmatrix} -1.6550 \\ 2.5280 \end{bmatrix}, \Sigma_2 = \begin{bmatrix} 1.9302 & -0.8914 \\ -0.8914 & 2.6287 \end{bmatrix}, \mu_3 = \begin{bmatrix} 4.0812 \\ 4.3114 \end{bmatrix}, \Sigma_3 = \begin{bmatrix} 7.7501 & 3.7525 \\ 3.7525 & 10.3827 \end{bmatrix}$$

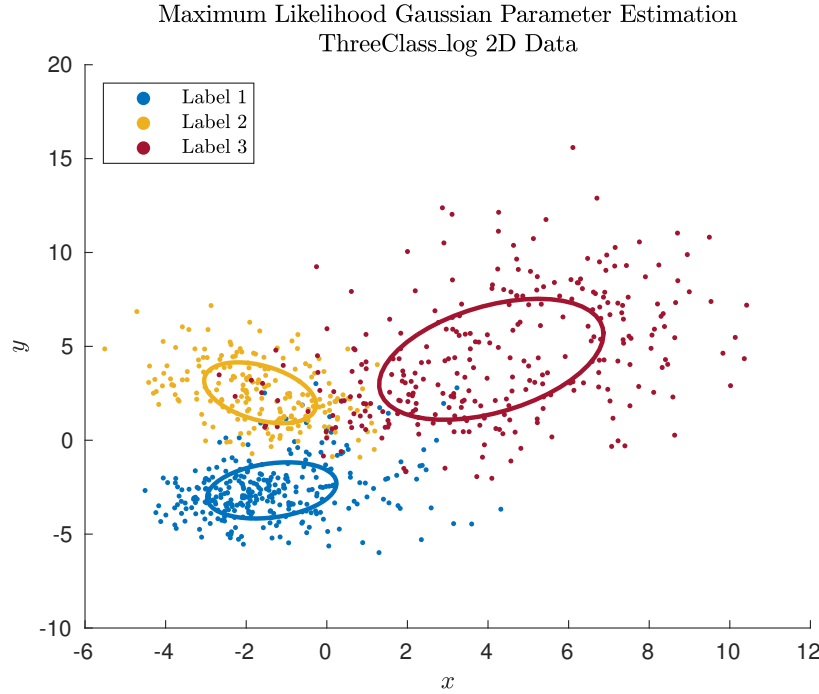


Figure 2: Estimated distribution using ML for model (a).

## Maximum Likelihood for Model (b)

Since we do not know  $z_i$ , we must marginalize over  $z$  to obtain the marginal log-likelihood of  $y_{1:N}$ . The resulting log-likelihood is therefore the incomplete log-likelihood. We cannot efficiently optimize this function directly as it is a sum of a log of a sum where the  $p(y_i|z_i = m; \theta_y)$  distribution is coupled to the unknown distribution of the R.V.  $Z_i$ . Therefore we have opted to use the expected complete log-likelihood to perform Expectation Maximization.

$$\begin{aligned}
L(y_{1:N}; \theta_j, \theta_y) &= p(y_{1:N} | L_i = j; \theta_y) = \prod_{i=1}^N \sum_{m=1}^2 p(L_i = j; \theta_L) p(z_i = m | L_i = j; \theta_y) p(y_i | z_i = m; \theta_y) \\
l(y_{1:N}; \theta_j, \theta_y) &= \log \left[ \prod_{i=1}^N \sum_{m=1}^2 p(L_i = j; \theta_L) p(z_i = m | L_i = j; \theta_y) p(y_i | z_i = m; \theta_y) \right] \\
l(y_{1:N}; \theta_j, \theta_y) &= \log[p(L_i = j; \theta_L)] + \sum_{i=1}^N \log \left[ \sum_{m=1}^2 p(L_i = j; \theta_L) p(z_i = m | L_i = j; \theta_y) p(y_i | z_i = m; \theta_y) \right]
\end{aligned}$$

For expectation maximization, we infer  $p(z_i = m | y_i; \theta_z, \theta_y)$  via the following:

$$\begin{aligned}
p(z_i = m | y_i; \theta_z, \theta_y) &= \frac{p(z_i = m; \theta_z, \theta_y) p(y_i | z_i; \theta_z, \theta_y)}{\sum_{n=1}^2 p(z_i = n; \theta_z) p(z_i = n | y_i; \theta_z, \theta_y)} \\
&= \frac{w_{mj} p(y_i | z_i = m; \theta_y)}{\sum_{n=1}^2 w_{nj} p(y_i | z_i = n)} \\
&= \frac{w_{mj} \mathcal{N}(y_i; \mu_m, \Sigma_m)}{\sum_{n=1}^2 w_{nj} \mathcal{N}(y_i; \mu_n, \Sigma_n)} \\
&= \gamma(z_{im})
\end{aligned}$$

$\gamma$  represents the responsibility, or the probability that data  $y_i$  within the label set  $j$  belongs to the cluster / component  $z_m$ . The calculation of  $\gamma$  is used as the expectation step for the EM algorithm. It is then used to update our parameters for the  $j^{th}$  class via the following equations:

$$\begin{aligned}
\hat{\mu}_{mj}^{\text{new}} &= \frac{1}{N_m} \sum_{i=1}^N \gamma(z_{im}) y_i \\
\hat{\Sigma}_{mj}^{\text{new}} &= \frac{1}{N_m} \sum_{i=1}^N \gamma(z_{im}) (y_i - \hat{\mu}_{mj})(y_i - \hat{\mu}_{mj})^T \\
\hat{w}_{mj}^{\text{new}} &= \frac{\sum_{i=1}^N \gamma(z_{im})}{N} = \frac{N_m}{N}
\end{aligned}$$

Note: We can update the  $j^{th}$  parameters independently using the training data associated with class  $j$  because the ICLL functions are associated with the same one-hot encoding that sifts out all data not associated with the parameters of  $j$ .

### Estimated Parameters for Model (b):

Parameters for Label 1

$$\mu_{11} = \begin{bmatrix} -1.8687 \\ -2.9637 \end{bmatrix}, \Sigma_{11} = \begin{bmatrix} 1.3125 & 0.1612 \\ 0.1612 & 1.1954 \end{bmatrix}, w_2 = 0.8143 \quad \mu_{21} = \begin{bmatrix} 0.8447 \\ -1.3926 \end{bmatrix}, \Sigma_{21} = \begin{bmatrix} 2.1391 & -0.9447 \\ -0.9447 & 4.8213 \end{bmatrix}, w_2 = 0.1857$$

Parameters for Label 2

$$\mu_{12} = \begin{bmatrix} -1.5592 \\ 1.9603 \end{bmatrix}, \Sigma_{12} = \begin{bmatrix} 1.9298 & -0.7605 \\ -0.7605 & 1.6368 \end{bmatrix}, w_2 = 0.7698 \quad \mu_{22} = \begin{bmatrix} -1.9752 \\ 4.4272 \end{bmatrix}, \Sigma_{22} = \begin{bmatrix} 1.7986 & -0.5394 \\ -0.5394 & 1.2613 \end{bmatrix}, w_2 = 0.2302$$

Parameters for Label 3

$$\mu_{13} = \begin{bmatrix} 5.8549 \\ 6.3759 \end{bmatrix}, \Sigma_{13} = \begin{bmatrix} 4.5488 & -0.3473 \\ -0.3473 & 8.1516 \end{bmatrix}, w_2 = 0.5277 \quad \mu_{23} = \begin{bmatrix} 2.0991 \\ 2.0044 \end{bmatrix}, \Sigma_{23} = \begin{bmatrix} 3.8830 & -0.3309 \\ -0.3309 & 2.7911 \end{bmatrix}, w_2 = 0.4723$$

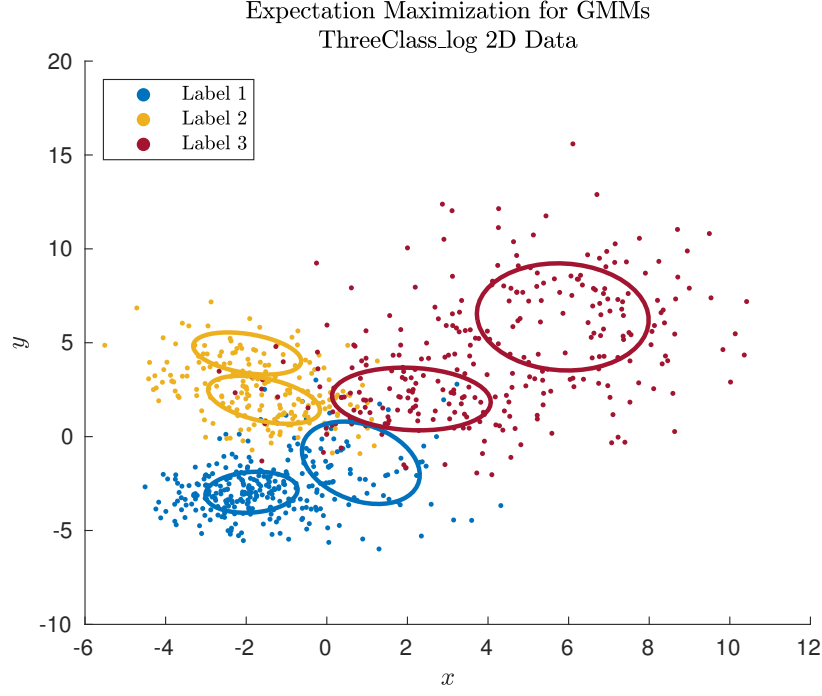


Figure 3: Estimated distribution using ML for model (a).

## Problem 2 - Baum-Welch for Supervisory Operator HMM

### Parameter Comparisons:

True	.25	.25	.25	.25
Estimated	.2078	.0453	.3376	.4093

10 Sequences

True	.25	.25	.25	.25
Estimated	.5314	.0762	.3210	.0714

50 Sequences

True	.25	.25	.25	.25
Estimated	0.3528	0.1352	.3090	.2029

100 Sequences

Table 1: Initial state probability  $p(x_0)$  estimates for 10, 50, and 100 observation sequences with 50 iterations.

True	.0200	.0190	0	.6660
	0	.0250	.5170	0
	.1630	.7690	.4660	0
	.8170	.1870	.0170	.3340
Estimated	.0897	.0252	0	.7669
	0	.0072	.4979	0
	.1586	.7843	.4832	0
	.7517	.1834	.0189	.2331

10 Sequences

True	.0200	.0190	0	.6660
	0	.0250	.5170	0
	.1630	.7690	.4660	0
	.8170	.1870	.0170	.3340
Estimated	.0546	.0235	0	.6971
	0	.0436	.5747	0
	.1672	.7588	.4120	0
	.7782	.1741	.0133	.3029

50 Sequences

True	.0200	.0190	0	.6660
	0	.0250	.5170	0
	.1630	.7690	.4660	0
	.8170	.1870	.0170	.3340
Estimated	.0686	.0317	0	.6801
	0	.0637	.5855	0
	.1681	.7336	.4033	0
	.7633	.1709	.0112	.3199

100 Sequences

Table 2: State transition probability  $p(x_k|x_{k-1})$  estimates for 10, 50, and 100 observation sequences with 50 iterations.

True	.0338	0	0	.3273
	.0934	0	0	.0949
	.1356	0	0	.0311
	.1031	0	0	.0125
	.1350	0	0	.0113
	.0289	0	0	.3354
	.0968	0	0	.1094
	.1409	0	0	.0488
	.1117	0	0	.0149
	.1208	0	0	.0144
	0	.0842	.7353	0
	0	.2048	.1869	0
	0	.2774	.0195	0
	0	.3336	.0267	0
	0	0	.0316	0
Estimated	0	0	0	.3796
	.0626	0	0	.0631
	.1405	0	0	.0080
	.1499	0	0	.0000
	.1090	0	0	.0000
	.0284	0	0	.3728
	.0775	0	0	.1029
	.1673	0	0	.0610
	.1163	0	0	.0055
	.1484	0	0	.0071
	0	.0637	.7721	0
	0	.1392	.1538	0
	0	.3087	.0581	0
	0	.4884	.0000	0
	0	0	.0160	0

10 Sequences

True	.0338	0	0	.3273
	.0934	0	0	.0949
	.1356	0	0	.0311
	.1031	0	0	.0125
	.1350	0	0	.0113
	.0289	0	0	.3354
	.0968	0	0	.1094
	.1409	0	0	.0488
	.1117	0	0	.0149
	.1208	0	0	.0144
	0	.0842	.7353	0
	0	.2048	.1869	0
	0	.2774	.0195	0
	0	.3336	.0267	0
	0	0	.0316	0
Estimated	.0002	0	0	.3501
	.0978	0	0	.0912
	.1212	0	0	.0186
	.1106	0	0	.0145
	.1325	0	0	.0091
	.0452	0	0	.3232
	.0925	0	0	.1201
	.1328	0	0	.0523
	.1313	0	0	.0122
	.1359	0	0	.0087
	0	.1210	.7743	0
	0	.2015	.1533	0
	0	.3600	.0222	0
	0	.3175	.0145	0
	0	0	.0356	0

50 Sequences

True	.0338	0	0	.3273
	.0934	0	0	.0949
	.1356	0	0	.0311
	.1031	0	0	.0125
	.1350	0	0	.0113
	.0289	0	0	.3354
	.0968	0	0	.1094
	.1409	0	0	.0488
	.1117	0	0	.0149
	.1208	0	0	.0144
	0	.0842	.7353	0
	0	.2048	.1869	0
	0	.2774	.0195	0
	0	.3336	.0267	0
	0	0	.0316	0
Estimated	.0216	0	0	.3321
	.0878	0	0	.0947
	.1387	0	0	.0157
	.1180	0	0	.0111
	.1437	0	0	.0087
	.0295	0	0	.3409
	.0794	0	0	.1258
	.1338	0	0	.0524
	.1087	0	0	.0142
	.1387	0	0	.0044
	0	.1324	.7527	0
	0	.2094	.1690	0
	0	.3289	.0254	0
	0	.3292	.0170	0
	0	0	.0358	0

100 Sequences

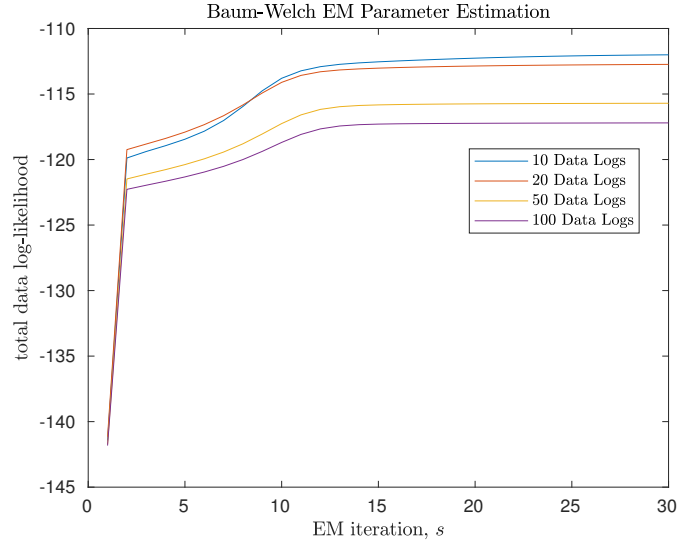
Table 3: Emission probability  $p(y_k|x_k)$  estimates for 10, 50, and 100 observation sequences with 50 iterations.**Log-likelihood Comparisons:**

Figure 4: Data log-likelihood average per iteration for 10, 20, 50, 100 observation sequences.

We initialized our parameters using an "informed" approach. The idea is that we might not know the exact values for each entry in our initial CPT's, but we know which values ought to be zero. We used a uniform distribution on all remaining non-zero entries appropriately to maintain row or column stochasticity. However, our results show that the lower count of observation sequences sometimes have a slightly higher log-likelihood. This is unexpected and perhaps may be due to a bug or incorrect implementation.

## Appendix - MATLAB Code

The following MATLAB code was used to generate the data presented above.

### hw2\_BaumWelch.m - main script

```
1 %% Header
2
3 % Filename:      hw2_BaumWelch.m
4 % Author:       Carl Stahoviak
5 % Date Created: 03/02/2019
6
7 clc;
8 clear;
9 close ALL;
10
11 %% Load Data
12
13 % load HMM CPT Parameters
14 load('nominal_hmm_params.mat')
15
16 %% Initialize CPT distributions for Baum-Welch
17
18 % load observation set for Baum-Welch implementation (y_obs)
19 % NOTE: each column of y_obs is an independent observation sequence
20 load('nominal_hmm_multi_logs.mat')
21
22 n = size(pxk_xkml,1); % number of states
23 T = size(y_obs,1);    % number of observations per data log
24 D = size(y_obs,2);    % number of unique data logs
25 p = size(pyk_xk,1);  % number of unique emission symbols
26
27 % number of unique data logs to use in E-step (Berkeley notation)
28 N_logs = 10;
29 log_sz = [10, 20, 50, 100];
30
31 % number of times M-step will be done
32 N_mstep = 30;
33
34 % total data log-likelihood per EM iteration
35 data_ll_total = zeros(N_mstep, size(log_sz, 2));
36
37 % total data log-likelihood per EM iteration
38 data_ll_mean = zeros(N_mstep, size(log_sz, 2));
39
40 % scatter plot size
41 sz = 5;
42
43
44 %% Baum-Welch - Attempt 2
45
46 % In this attempt, the e-step (resulting in the calculation of {el_nalpha,
47 % el_nbeta, elngamma}) will be run for all D data logs, and a single m-step
48 % will be done at the end of this. So rather than performing D number of
49 % e-step/m-step iterations (which failed at the second iteration), D number
50 % of e-step iterations will be performed and then a single m-step will be
51 % done to update the CPT parameter estimates.
52
53 figure(1)
54 ax1 = gca;
55 title('Baum-Welch EM Parameter Estimation', 'Interpreter', 'latex');
56 xlabel('EM iteration, $s$', 'Interpreter', 'latex');
57 ylabel('data log-likelihood', 'Interpreter', 'latex');
58 hold on;
59
60 for i=1:size(log_sz, 2)
```

```

61
62     N_logs = log.sz(i);
63
64     % initialize CPTs
65     trans_prob = zeros(n,n,N_mstep+1);      % row-stochastic
66     obs_prob   = zeros(p,n,N_mstep+1);      % column-stochastic (is this correct?)
67     init_distr = zeros(n,N_mstep+1);        % column-stochastic
68
69     % initialize Conditional Probability Tables (CPTs)
70     type = 'informed';
71     [ init_distr(:,1), trans_prob(:, :, 1), obs_prob(:, :, 1) ] = initCPTs( ...
72         pxk_xkml', pyk_xk, px0, type);
73
74     % data log-likelihood (for each data log per EM iteration)
75     data_ll = zeros(N_logs, N_mstep);
76
77     for s=1:N_mstep % number of times M-step will be done
78
79         fprintf('EM iteration: %d\n', s);
80
81         % init variables
82         eln_alpha = zeros(n,T+1,N_logs);    % T+1 columns because alpha(0) is computed
83         eln_beta  = zeros(n,T+1,N_logs);    % T+1 columns because beta(0) is computed
84         eln_gamma = zeros(n,T+1,N_logs);    % log-posterior
85         gamma     = zeros(n,T+1,N_logs);    % true posterior
86         eln_xi    = zeros(n,n,T,N_logs);    % log of probability xi(i,j,k)
87         xi        = zeros(n,n,T,N_logs);
88
89         for d=1:N_logs % iterate over D number of data logs
90             obs_seq = y_obs(:,d);
91
92             % do E-step for each data log
93             [eln_alpha(:, :, d), eln_beta(:, :, d), eln_gamma(:, :, d), ...
94                 gamma(:, :, d), eln_xi(:, :, :, d), xi(:, :, :, d)] = baumwelchEstep( ...
95                 init_distr(:,s), trans_prob(:, :, s), obs_prob(:, :, s), obs_seq);
96
97             % calculate data log-likelihood for the e-step given the current
98             % CPT parameter estimates
99             nonNaN_idx = ~isnan(eln_alpha(:,end,d));
100             data_ll(d,s) = eln(sum(exp(eln_alpha(nonNaN_idx,end,d))));
101         end
102
103         % do M-step given D iterations of the E-step
104         [init_distr(:,s+1), trans_prob(:, :, s+1), obs_prob(:, :, s+1) ] = ...
105             baumwelchMstep( p, eln_gamma, gamma, eln_xi, xi, y_obs);
106
107         % plot the data log-likelihood of each data log at EM iteration s
108         scatter(ax1,s*ones(N_logs,1),data_ll(:,s),sz,'filled');
109
110         % compute the total data log-likelihood per EM iteration
111         data_ll_total(s,i) = sum(data_ll(:,s));
112         data_ll_mean(s,i) = mean(data_ll(:,s));
113     end
114 end
115
116
117 fprintf('\nFinal (sum) data log-likelihood =\n')
118 disp([log.sz; data_ll_total(end,:)]);
119 fprintf('Final (avg) data log-likelihood =\n\n')
120 disp([log.sz; data_ll_mean(end,:)]);
121
122 fprintf('Estimated Initial Distribution\n')
123 disp(init_distr(:,end));
124 disp(sum(init_distr(:,end),1));
125
126 fprintf('True State Transition CPT\n')
127 disp(pxk_xkml);
128 fprintf('%d Estimated State Transition CPT\n', s)

```



```

129 disp(trans_prob(:, :, end));
130 disp(1-sum(trans_prob(:, :, end), 1));
131
132 fprintf('True Emission CPT\n')
133 disp(pyk_xk);
134 fprintf('Estimated Emission CPT\n')
135 disp(obs_prob(:, :, end));
136 disp(1-sum(obs_prob(:, :, end), 1));
137
138 figure(2)
139 for i=1:size(log_sz, 2)
140     plot(data_all_mean(:, i));
141     title('Baum-Welch EM Parameter Estimation', 'Interpreter', 'latex');
142     xlabel('EM iteration, $s$', 'Interpreter', 'latex');
143     ylabel('total data log-likelihood', 'Interpreter', 'latex');
144     hold on;
145     hdl = legend();
146 end
147 set(hdl, 'Interpreter', 'latex');

```

### initCPTs.m - Initialize CPT distributions

```

1 function [ init_distr, trans_prob, obs_prob ] = ...
2     initCPTs( pxk_xkml, pyk_xk, px0, type )
3 %UNTITLED Summary of this function goes here
4 % Detailed explanation goes here
5
6 % NOTE: in accordance with Rabiner/Mann notation, the trans_prob matrix
7 % must be a row-stochastic matrix
8
9 n = size(pxk_xkml, 1); % number of states
10 p = size(pyk_xk, 1); % number of unique emission symbols
11
12 if strcmp(type, 'truth')
13     init_distr = px0;
14     trans_prob = pxk_xkml; % row-stochastic
15     obs_prob = pyk_xk;
16
17 elseif strcmp(type, 'informed')
18     % "informed" uniform initialization
19     % initialize CPTs WITH knowledge of where zeros exist in the tables
20     init_distr = (1/n)*ones(n, 1);
21
22     trans_prob = [1/3 0 1/3 1/3;
23                 1/4 1/4 1/4 1/4;
24                 0 1/3 1/3 1/3;
25                 1/2 0 0 1/2];
26
27     obs_prob = [(1/10)*ones(10, 1), zeros(10, 2), (1/10)*ones(10, 1);
28                zeros(4, 1), (1/4)*ones(4, 1), (1/5)*ones(4, 1), zeros(4, 1);
29                0, 0, 1/5, 0];
30
31 elseif strcmp(type, 'uninformed')
32     % "un-informed" uniform initialization
33     % initialize CPTs WITHOUT knowledge of where zeros exist in the tables
34     init_distr = (1/n)*ones(n, 1);
35     trans_prob = (1/n)*ones(n);
36     obs_prob = (1/p)*ones(p, n);
37
38 elseif strcmp(type, 'rand')
39     init_distr = rand(n, 1);
40     init_distr = init_distr./sum(init_distr); % normalize
41
42     trans_prob = rand(n);
43     for i=1:n
44         % normalize rows
45         trans_prob(i, :) = trans_prob(i, :)./sum(trans_prob(i, :));

```

```

46     end
47
48     obs_prob = rand(p,n);
49     for i=1:n
50         % normalize columns
51         obs_prob(:,i) = obs_prob(:,i)./sum(obs_prob(:,i));
52     end
53
54 else
55     error('Improperly specified initialization type')
56 end
57
58 end

```

### baumwelch\_Estep.m - Baum-Welch E-Step

```

1 function [ eln_alpha, eln_beta, eln_gamma, gamma, eln_xi, xi ] = ...
2     baumwelch_Estep( init_distr, trans_prob, obs_prob, y_obs )
3
4 % Mann log-weighted (numerically-stable) forward-backward alg.
5 eln_alpha = forward_eln( init_distr, trans_prob, obs_prob, y_obs );
6 eln_beta = backward_eln( trans_prob, obs_prob, y_obs );
7
8 % get the log-posterior and posterior distributions (exact inference)
9 % [ eln_gamma, gamma ] = posterior_elnfb( eln_alpha(:,2:end), eln_beta );
10 [ eln_gamma, gamma ] = posterior_elnfb( eln_alpha, eln_beta );
11
12 % get log of probability xi(i,j,k)
13 [ eln_xi, xi ] = eln_xi( eln_alpha, eln_beta, ...
14     trans_prob, obs_prob, y_obs );
15
16 end

```

### forward\_eln.m - extended-logarithm forward pass

```

1 function [ eln_alpha ] = forward_eln( px0, trans_prob, obs_prob, y_obs )
2
3 % NOTE: in accordance with Rabiner/Mann notation, the trans_prob matrix
4 % must be a row-stochastic matrix
5
6 n = size(px0,1);
7 T = size(y_obs,1);
8 eln_alpha = zeros(n,T+1);
9
10 % initialization
11 for i=1:n
12     eln_alpha(i,1) = eln(px0(i,1));
13 end
14
15 for k=2:T+1
16     for j=1:n
17         logalpha = NaN;
18         for i=1:n
19             logalpha = elnsum( logalpha, ...
20                 elnprod( eln_alpha(i,k-1), eln(trans_prob(i,j)) ));
21         end
22         eln_alpha(j,k) = elnprod( logalpha, ...
23             eln(obs_prob(y_obs(k-1),j)) );
24     end
25 end
26
27 end

```

### backward\_eln.m - extended-logarithm backward pass

```

1 function [ eln_beta ] = backward_eln( trans_prob, obs_prob, y_obs )
2
3 % NOTE: in accordance with Rabiner/Mann notation, the trans_prob matrix
4 % must be a row-stochastic matrix
5
6 n = size(trans_prob,1);
7 T = size(y_obs,1);
8
9 % initialization
10 eln_beta = zeros(n,T+1);
11
12 for k=T:-1:1
13     for i=1:n
14         logbeta = NaN;
15         for j=1:n
16             logbeta = elnsum( logbeta, ...
17                             elnprod( eln(trans_prob(i,j)), ...
18                                     elnprod( eln(obs_prob(y_obs(k),j)), eln_beta(j,k+1) )));
19         end
20         eln_beta(i,k) = logbeta;
21     end
22 end

```

**elnfb\_posterior.m** - extended-logarithm posterior calculation

```

1 function [ eln_gamma, gamma ] = posterior_elnfb( eln_alpha, eln_beta )
2
3 if size(eln_alpha,2) ≠ size(eln_beta,2)
4     error('eln_alpha, eln_beta size mismatch')
5
6 else
7     n = size(eln_alpha,1);
8     T = size(eln_alpha,2);
9
10    % initialization
11    eln_gamma = zeros(n,T);    % log-posterior
12    gamma = zeros(n,T);      % true posterior
13
14    for k=1:T
15        normalizer = NaN;
16        for i=1:n
17            eln_gamma(i,k) = elnprod(eln_alpha(i,k),eln_beta(i,k));
18            normalizer = elnsum(normalizer,eln_gamma(i,k));
19        end
20
21        for i=1:n
22            eln_gamma(i,k) = elnprod(eln_gamma(i,k),-normalizer);
23            gamma(i,k) = exp(eln_gamma(i,k));
24        end
25        % sanity check
26        fprintf('1 - sum(gamma(:,%d)) = %e\n', k, ...
27              1-nansum(exp(eln_gamma(:,k))))
28    end
29 end

```

**elnxi.m** - extended-logarithm xi calculation

```

1 function [ eln_xi, xi ] = elnxi( eln_alpha, eln_beta, trans_prob, obs_prob, y_obs )
2 %ELNXI - Computes the log of xi(i,j,k)
3 % xi(i,j,k) is the probability of being in state x_i and timestep k, and
4 % state x_j at timestep k+1 given a model lambda and observation sequence
5 % O (in Rabiner notation)
6
7 % NOTE: in accordance with Rabiner/Mann notation, the trans_prob matrix
8 % must be a row-stochastic matrix

```

```

9
10 n = size(trans_prob,1);
11 T = size(y_obs,1);
12
13 % initialization
14 eln_xi = zeros(n,n,T);
15 xi = zeros(n,n,T);
16
17 for k=1:T
18     normalizer = NaN;
19     for i=1:n
20         for j=1:n
21             eln_xi(i,j,k) = elnprod( eln_alpha(i,k), ...
22                                     elnprod( eln(trans_prob(i,j)), ...
23                                               elnprod( eln(obs_prob(y_obs(k),j)), eln_beta(j,k+1) )));
24             normalizer = elnsum( normalizer, eln_xi(i,j,k) );
25         end
26     end
27
28     for i=1:n
29         for j=1:n
30             eln_xi(i,j,k) = elnprod( eln_xi(i,j,k), -normalizer);
31             xi(i,j,k) = exp(eln_xi(i,j,k));
32         end
33     end
34 end
35
36 end

```

#### baumwelch\_Mstep.m - Baum-Welch M-Step

```

1 function [ init_distr_hat, trans_prob_hat, obs_prob_hat ] = ...
2     baumwelchMstep( p, eln_gamma, gamma, eln_xi, xi, y_obs )
3 %UNTITLED2 Summary of this function goes here
4 % Detailed explanation goes here
5
6 if (size(gamma,2) ≠ size(xi,3)+1) && (size(gamma,3) ≠ size(xi,4))
7     error('gamma, xi size mismatch')
8
9 else
10     % Implement equations from Berkeley paper
11     % init_distr_hat = updateInitDistr( gamma );
12     % trans_prob_hat = updateTransitionProb( gamma, xi);
13     % obs_prob_hat = updateEmissionProb( p, gamma, y_obs );
14
15     % Implement extended-log versions of the parameter updates
16     % from the Berkeley paper:
17     % WORKING!! - Now need to update functions above to achieve similar
18     % results!
19     init_distr_hat = updateInitDistr.eln( eln_gamma );
20     trans_prob_hat = updateTransitionProb.eln( eln_gamma, eln_xi);
21     obs_prob_hat = updateEmissionProb.eln( p, eln_gamma, y_obs );
22 end

```

#### updateInitDistr.eln.m - Baum-Welch Initial Distribution Update

```

1 function [ init_distr_hat ] = updateInitDistr.eln( eln_gamma )
2
3 n = size(eln_gamma,1);
4 D = size(eln_gamma,3);
5
6 init_distr_hat = zeros(n,1);
7
8 % get initial distribution estimate
9 for d=1:D
10     for i=1:n

```

```

11         init_distr_hat(i,1) = init_distr_hat(i,1) + exp(eln_gamma(i,1,d));
12     end
13 end
14
15 init_distr_hat = init_distr_hat/D;
16
17 % normalize
18 % init_distr_hat = init_distr_hat./sum(init_distr_hat);
19
20 end

```

### updateTransitionProb\_elm.m - Baum-Welch Transition Probability Update

```

1 function [ trans_prob_hat ] = updateTransitionProb_elm( elm_gamma, elm_xi)
2
3 n = size(elm_gamma,1);
4 T = size(elm_gamma,2) - 1;
5 D = size(elm_gamma,3);
6
7 trans_prob_hat = zeros(n,n);
8
9 % get transition probability estimate
10 for i=1:n
11     for j=1:n
12         numerator = NaN;
13         denominator = NaN;
14
15         for d=1:D
16             for k=2:T+1
17                 numerator = elnsum(numerator, elm_xi(i,j,k-1,d));
18                 denominator = elnsum(denominator, elm_gamma(i,k-1,d));
19             end % end k
20         end % end d
21
22         trans_prob_hat(i,j) = exp(elnprod(numerator, -denominator));
23     end % end j
24 end % end i
25
26 end

```

### updateEmissionProb\_elm.m - Baum-Welch Emission Probability Update

```

1 function [ obs_prob_hat ] = updateEmissionProb_elm( p, elm_gamma, y_obs )
2 %UNTITLED5 Summary of this function goes here
3 % Detailed explanation goes here
4
5 n = size(elm_gamma,1);
6 T = size(elm_gamma,2) - 1;
7 D = size(elm_gamma,3);
8
9 obs_prob_hat = zeros(p,n);
10
11 % get emission probability estimate
12 for j=1:p
13     for i=1:n
14         numerator = NaN;
15         denominator = NaN;
16
17         for d=1:D
18             for k=2:T+1
19                 %
20                 % symbol j observed at time step k in data log d
21                 if y_obs(k-1,d) == j
22                     numerator = elnsum(numerator, elm_gamma(i,k,d));
23                 end
24                 denominator = elnsum(denominator, elm_gamma(i,k,d));

```

```

25         end % end k
26     end % end d
27
28     obs_prob_hat(j,i) = exp(elnprod(numerator, -denominator));
29 end % end i
30 end % end j
31
32 end

```

## hw2\_MLE.m - Maximum Likelihood and Expectation Maximization

```

1  %% Header
2
3  % Filename:      hw2_MLE.m
4  % Author:       Carl Stahoviak
5  % Date Created: 03/02/2019
6
7  clc;
8  clear;
9  close ALL;
10
11 %% Load Data
12
13 % load MLE Data Log
14 load('ThreeClass_log.mat')
15 y_obs = ThreeClass_log;
16
17 %% Plot Data
18
19 colors = [0,      0.4470, 0.7410;
20           0.9290, 0.6940, 0.1250;
21           0.6350, 0.0780, 0.1840];
22
23 figure(1);
24 ax(1) = gca;
25 title({'Maximum Likelihood Gaussian Parameter Estimation', ...
26        'ThreeClass_log 2D Data'}, 'Interpreter', 'latex');
27 xlabel('$x$', 'Interpreter', 'latex');
28 ylabel('$y$', 'Interpreter', 'latex');
29 hold on;
30
31 for i=1:3
32     figure(i+1);
33     ax(i+1) = gca;
34     title({'Expectation Maximization for GMMs', ...
35            strcat('Data Label', '$\_$', num2str(i))}, 'Interpreter', 'latex');
36     xlabel('$x$', 'Interpreter', 'latex');
37     ylabel('$y$', 'Interpreter', 'latex');
38     hold on;
39     pause on;
40 end
41
42 figure(5);
43 ax(5) = gca;
44 title({'Expectation Maximization for GMMs', ...
45        'ThreeClass_log 2D Data'}, 'Interpreter', 'latex');
46 xlabel('$x$', 'Interpreter', 'latex');
47 ylabel('$y$', 'Interpreter', 'latex');
48 hold on;
49 pause on;
50
51 figure(6);
52 ax(6) = gca;
53 title({'Expectation Maximization for GMMs', ...
54        'Unknown Data Labels'}, 'Interpreter', 'latex');
55 xlabel('$x$', 'Interpreter', 'latex');
56 ylabel('$y$', 'Interpreter', 'latex');

```

```

57 hold on;
58 pause on;
59
60
61 %% Maximum Likelihood Estimation - Multivariate Gaussian
62
63 % For Model A, the labels {pi-j = 1,2,3} associated with each data point
64 % are known. In this case, it is straightforward to compute the gaussian
65 % parameters theta = {mu-j, sigma-j} = {mu1, sigma1, mu2, sigma2, mu3,
66 % sigma3} via Maximum Likelihood Estimation (MLE).
67
68 % get one-pass ML estimates of mean and variance of each distribution
69
70 mu = zeros(2,max(y_obs(:,1)));
71 sigma = zeros(2,2,max(y_obs(:,1)));
72
73 idx_one = (y_obs(:,1) == 1);
74 idx_two = (y_obs(:,1) == 2);
75 idx_three = (y_obs(:,1) == 3);
76
77 for i=1:length(ax)
78     if (i == 1) || (i == 5)
79         h1 = scatter(ax(i),y_obs(idx_one,2),y_obs(idx_one,3), ...
80                     5, colors(y_obs(idx_one,1,:),:),'filled');
81
82         h2 = scatter(ax(i),y_obs(idx_two,2),y_obs(idx_two,3), ...
83                     5, colors(y_obs(idx_two,1,:),:),'filled');
84
85         h3 = scatter(ax(i),y_obs(idx_three,2),y_obs(idx_three,3), ...
86                     5, colors(y_obs(idx_three,1,:),:),'filled');
87
88         hdl = legend(ax(i),[h1(1), h2(1), h3(1)],'Label 1','Label 2','Label 3');
89         set(hdl,'Interpreter','latex','Location','Northwest','Autoupdate','off');
90     end
91 end
92
93 [mu(:,1),sigma(:,:,1)] = gaussian_mle( y_obs(idx_one,2:3)' );
94 plot_gaussian_ellipsoid(mu(:,1), sigma(:,:,1), colors(1,:), 1, [], ax(1));
95
96 [mu(:,2),sigma(:,:,2)] = gaussian_mle( y_obs(idx_two,2:3)' );
97 plot_gaussian_ellipsoid(mu(:,2), sigma(:,:,2), colors(2,:), 1, [], ax(1));
98
99 [mu(:,3),sigma(:,:,3)] = gaussian_mle( y_obs(idx_three,2:3)' );
100 plot_gaussian_ellipsoid(mu(:,3), sigma(:,:,3), colors(3,:), 1, [], ax(1));
101
102 %% Expectation Maximization - Gaussian Mixture Models (GMMs)
103
104 % For Model B, the data labels {pi-j = 1,2,3} are known, but the variable
105 % Z-i, the sub-distribution labels {m = 1,2} are unknown (latent
106 % variables), and we can use Expectation Maximization (EM) to learn the
107 % parameters of the GMM - {w1-j, w2-j, mu1-j, sigma1-j, mu2-j, sigma2-j}.
108
109 % Since the data labels {pi-j = 1,2,3} are known, we can use EM to
110 % determine the parameter set {w1-j, w2-j, mu1-j, sigma1-j, mu2-j, sigma2-j}
111 % associated with each label j. This process for a given label j is
112 % independent of every other label j.
113
114 n = 2; % dimensionality of data
115 M = 2; % number of sub-distributions within label j
116
117 % number of times EM algorithm will be done
118 N_iter = 100;
119
120 for j=1:3
121
122     % get data with label j
123     idx = (y_obs(:,1) == j);
124     X = y_obs(idx,2:3)';

```

```

125 Npoints = size(X,2);
126
127 % init parameters
128 mu      = zeros(n,Niter,M);
129 sigma    = zeros(n,n,Niter,M);
130 weights  = zeros(M,Niter);
131 r        = zeros(M,Npoints,Niter);
132
133 % get intial parameter estimates
134 lower = [min(X(1,:)), min(X(2,:))]; % min [x,y]
135 upper = [max(X(1,:)), max(X(2,:))]; % max [x,y]
136 for m=1:M
137     for i=1:n
138         mu(i,1,m) = (upper(i)-lower(i)).*rand(1) + lower(i);
139     end
140     sigma(:, :, 1, m) = [5, 0.5; 0.5 5];
141 end
142 weights(:,1) = ones(M,1)/M;
143
144 for s=1:Niter
145     % clear axes and replot data
146     cla(ax(j+1));
147     scatter(ax(j+1), y_obs(idx,2), y_obs(idx,3), 5, ...
148             colors(j,:), 'filled');
149
150     % Gaussian EM - E-step
151     [ r(:, :, s) ] = gaussian_em_estep( X, squeeze(mu(:,s,:)), ...
152                                         squeeze(sigma(:, :, s, :)), weights(:,s) );
153
154     % Gaussian EM - M-step
155     [ mu(:,s+1,:), sigma(:, :, s+1,:), weights(:,s+1) ] = ...
156         gaussian_em_mstep( X, r(:, :, s), squeeze(mu(:,s,:)) );
157
158     for m=1:M
159         % plot gaussian ellipsoid m
160         plot_gaussian_ellipsoid(mu(:,s+1,m), sigma(:, :, s+1,m), ...
161                                 colors(j,:), 1, [], ax(j+1));
162     end
163     pause(.1);
164
165 end
166
167 % plot m sub-distributions within label j
168 for m=1:M
169     % plot gaussian ellipsoid m
170     plot_gaussian_ellipsoid(mu(:,s+1,m), sigma(:, :, s+1,m), ...
171                             colors(j,:), 1, [], ax(5));
172 end
173
174 fprintf('data label %d:\n', j);
175 disp([mu(:,end,1), mu(:,end,2)])
176 disp(sigma(:, :, end, 1))
177 disp(sigma(:, :, end, 2))
178 disp(weights(:,end));
179
180 end

```

### gaussian\_mle.m - Maximum Likelihood Estimators

```

1 function [ mu, sigma ] = gaussian_mle( X )
2
3 n = size(X,1); % number of variables in multivariate Gaussian
4 M = size(X,2); % number of observations of each n-dim realization of X
5
6 mu = sum(X,2)/M;
7
8 sigma = zeros(n);

```



```

9  for i=1:size(X,2)
10     sigma = sigma + (X(:,i) - mu)*(X(:,i) - mu)';
11  end
12
13  sigma = sigma/M;

```

### gaussian\_em\_estep.m - Expectation Maximization E-Step

```

1  function [ r ] = gaussian_em_estep( X, mu, sigma, weights )
2
3  M = size(mu,2);           % number of sub-distributions within label j
4  Npoints = size(X,2);      % number of data points in label j
5
6  % fprintf('estep: [M, Npoints] = [%d, %d]\n', M, Npoints);
7
8  % init responsibility
9  r = zeros(M,Npoints);
10
11  % Gaussian EM - E-step
12  for m = 1:M
13      for i = 1:Npoints
14          numerator = weights(m,1) * ...
15              mvnpdf(X(:,i),mu(:,m),sigma(:, :, m));
16
17          denominator = 0;
18          for j=1:M
19              denominator = denominator + weights(j,1) * ...
20                  mvnpdf(X(:,i),mu(:,j),sigma(:, :, j));
21          end
22
23          % compute "responsibility" that data point i belongs to cluster m
24          r(m,i) = numerator/denominator;
25      end
26  end
27
28  end

```

### gaussian\_em\_mstep.m - Expectation Maximization M-Step

```

1  function [ mu, sigma, weights ] = gaussian_em_mstep( X, r, mu )
2
3  n = size(mu,1);           % dimensionality of data
4  M = size(mu,2);           % number of sub-distributions within label j
5  Npoints = size(X,2);      % number of data points in label j
6
7  % fprintf('mstep: [M, Npoints] = [%d, %d]\n', M, Npoints);
8
9  % initialize
10  mu_new = zeros(n,M);
11  sigma_new = zeros(n,n,M);
12  weights_new = zeros(M,1);
13
14  % Gaussian EM - M-step
15  for m=1:M
16      for i=1:Npoints
17          % update mean of distribution m
18          mu_new(:,m) = mu_new(:,m) + r(m,i)*X(:,i);
19
20          % update covariance matrix of distribution m
21          sigma_new(:, :, m) = sigma_new(:, :, m) + r(m,i) * ...
22              (X(:,i) - mu(:,m))*(X(:,i) - mu(:,m))';
23
24          % update weight of distribution m
25          weights_new(m,1) = weights_new(m,1) + r(m,i);
26      end
27  end

```

```

28     % sanity check - do these match?
29     fprintf('Nm(%d) = %f\n', m, weights_new(m,1))
30     fprintf('Nm(%d) = sum(r(m,:)) = %f\n\n', m, sum(r(m,:)))
31
32     % current number of points assigned to cluster m
33     Nm = sum(r(m,:));
34
35     % normalize by total number of data points
36     mu_new(:,m) = mu_new(:,m)/Nm;
37     sigma_new(:, :,m) = sigma_new(:, :,m)/Nm;
38     weights_new(m,1) = weights_new(m,1)/Npoints;
39
40 end
41
42     % resize parameters
43     mu = reshape(mu_new,n,1,M);
44     sigma = reshape(sigma_new,n,n,1,M);
45     weights = weights_new;
46
47 end

```

## Mann Extended-Logarithm Functions Library

```

1 function [ out ] = eexp( x )
2 % EEXP - Extended Exponential function
3
4     % if x == 'LOGZERO'
5     if isnan(x)
6         out = 0;
7     else
8         out = exp(x);
9     end
10
11 end
12
13 function [ out ] = eln( x )
14 % ELN - Extended Natural Logarithm function
15 % Computes the extended natural logarithm as defined by Mann, 2006
16
17 if x == 0
18     % out = 'LOGZERO';
19     out = NaN;
20 elseif x > 0
21     out = log(x);
22 else
23     error('eln() negative input error')
24 end
25
26 end
27
28 function [ prod ] = elnprod( eln_x, eln_y )
29 % ELNSUM - Extended Logarithm Product function
30 % Computes the extended logarithm of the product of x and y given as
31 % given as inputs the extended logarithm of x and y, as defined by Mann, 2006
32
33 % if strcmp(eln(x),'LOGZERO') || strcmp(eln(y),'LOGZERO')
34 if isnan(eln_x) || isnan(eln_y)
35     % prod = 'LOGZERO';
36     prod = NaN;
37 else
38     prod = eln_x + eln_y;
39 end
40
41 end
42
43 function [ sum ] = elnsum( eln_x, eln_y )
44 % ELNSUM - Extended Logarithm Sum function

```

```

45 % Computes the extended logarithm of the sum of x and y given as inputs
46 % the extended logarithm of x and y, as defined by Mann, 2006
47
48 % if strcmp(eLN(x),'LOGZERO') || strcmp(eLN(y),'LOGZERO')
49 if isnan(eLN.x) || isnan(eLN.y)
50     % if strcmp(eLN(x),'LOGZERO')
51     if isnan(eLN.x)
52         sum = eLN.y;
53     else
54         sum = eLN.x;
55     end
56 else
57     if eLN.x > eLN.y
58         sum = eLN.x + eLN(1 + exp(eLN.y-eLN.x));
59     else
60         sum = eLN.y + eLN(1 + exp(eLN.x-eLN.y));
61     end
62 end
63
64 end

```